

Yeah, my bad — you're right to call that out. I just re-opened the **actual zip** and here's what's *really* going on, end-to-end, based on the code that's in there.

## What this zip actually contains

It's a **capsule format + (demo) pipeline harness** for one concrete VM/AIR ("geom") that shows how you want capsules to be built and verified.

What's included:

- `scripts/run_geom_pipeline.py` — **producer**: runs Geometry → STC → STARK proof → (optional) Nova recursion → writes a `bef_capsule_v1` JSON/BIN capsule + an artifact manifest + row archive.
- `scripts/verify_capsule.py` — **verifier**: "accept/reject with stable error codes", verifying *hashes, statement-binding, proof validity, policy registry membership, signature/ACL, and DA sampling audit*.
- `bef_zk/spec/{trace.py, statement.py}` — defines your two key *adoption surface primitives*:
  - `TraceSpecV1` + `trace_spec_hash`
  - `StatementV1` + `statement_hash`
- `trace_statement_spec.md` + `README.md` — explains the intention.

What's **not** included (but referenced heavily): the actual `bef_zk` implementation (codec, DA providers, STC VC, AIR, prover/verifier, error codes, etc.). So the zip is a "*capsule interface + geom demo driver*" snapshot, not the whole repo.

---

## The real architecture (as implemented)

### 1) Producer pipeline: `run_geom_pipeline.py`

This script does five big things:

#### A) Simulate a trace

- `simulate_trace(GEOM_PROGRAM, params, init)` produces rows.
- It also derives a **column schema**, hashes it, and sticks `sha256 : <schema_hash>` into `record_schema_ref`.

#### B) Create TraceSpecV1

```

TraceSpecV1(
    spec_version="1.0",
    trace_format_id="GEOM_AIR_V1",
    record_schema_ref="sha256:<schema_hash>",
    encoding_id=<selected_encoding>,
    field_modulus_id="goldilocks_61",
)
trace_spec_hash = compute_trace_spec_hash(trace_spec)

```

So: the trace isn't "just bytes." It's *typed* by a content-addressed schema reference.

### C) Prove (STARK) with statement binding

During proving, you build a `StatementV1` inside `_build_statement_hash(...)`, and critically:

- `policy_hash` = `sha256(policy_file_bytes)`
- `trace_root` = `row_commitment.params["root"]` (Merkle root over chunk roots)
- `public_inputs` are concrete geom values (`final_m11`, `final_m12`, etc.)
- `anchors` optionally include an external anchor ref (L1 block hash, dataset hash, whatever)

Then `statement_hash` is computed and **fed into the transcript** via `statement_hash_fn`.

That's the "no proof reuse under a different meaning" property you want.

### D) Export STC row archive + chunk roots

- It extracts `chunk_handles` and `chunk_roots_hex` from proof metadata, then persists:
  - `row_archive/chunk_roots.json`
  - `row_archive/chunk_roots.bin`
  - plus `chunk_roots_digest` (sha256 over the `.bin`)
- It constructs:
  - `chunk_meta` (num\_chunks, chunk\_len, sizes)
  - `row_index_ref` (commitment root, arity, pointer to archive)

## E) Emit a capsule

It emits `bef_capsule_v1` containing:

- `trace_spec` and `trace_spec_hash`
- `statement` and `statement_hash`
- `policy: {policy_id, policy_version, policy_path, policy_hash}`
- `da_policy:`
  - `verification_level: "probabilistic_da_sampling"`
  - provider = local filesystem archive root
  - `k_samples` defaults to `ceil(log(1/epsilon)/delta)` with hardcoded (`delta=0.1, epsilon=1e-6`)
- `proofs.geom:`
  - proof paths, sizes
  - `row_archive` metadata
  - `formats` (json/bin with `sha256_payload_hash` per format)
- optional `proofs.nova` block and `trace_commitment = nova_state`

Finally:

- computes `capsule_hash = compute_capsule_hash(capsule_without_capsule_hash)`
  - optionally adds `authorship` signature (secp256k1 recoverable sig over capsule hash)
- 

## 2) Verifier pipeline: `verify_capsule.py`

This is honestly the most “product-like” piece in the zip. It’s a strict accept/reject verifier with stable error codes.

It verifies in layers:

### (1) Parse + schema check

- Must be dict
- `schema == "bef_capsule_v1"`

### (2) Capsule hash integrity

- Requires `capsule_hash` field.
- Recomputes canonical hash after removing `capsule_hash`.
- Mismatch ⇒ reject.

### (3) TraceSpec integrity

- Requires `trace_spec` object AND `trace_spec_hash`.
- Recomputes `compute_trace_spec_hash(TraceSpecV1.from_obj(...))`
- Any mismatch  $\Rightarrow$  `E053_PROOF_STATEMENT_MISMATCH`

### (4) Optional policy registry pinning

If you provide `--policy`, `--policy-inclusion-proof`, and `--policy-registry-root`:

- checks `sha256(policy_file) == capsule.policy.policy_hash`
- verifies a Merkle proof that the policy hash is included under the trusted registry root

So you support both “self-asserted policy hash” and “policy hash pinned to a governance root.”

### (5) Proof artifact integrity

- Locates proof path from capsule.
- If the capsule includes `sha256_payload_hash` for that proof format, it recomputes payload hash (with a header-skip rule) and checks it.

### (6) Binding: proof row commitment must match capsule metadata

This is that anti-tamper invariant you were talking about:

- compares root / chunk\_len / arity between:
  - proof’s internal `row_commitment` params
  - capsule’s `chunk_meta + row_index_ref`

Mismatch  $\Rightarrow$  `E053_PROOF_STATEMENT_MISMATCH`

### (7) Statement integrity + cross-binding

- Requires `statement` object and `statement_hash`.
- Recomputes statement hash, checks it matches.
- Also checks:
  - `statement.trace_spec_hash` equals `capsule.trace_spec_hash`
  - `statement.policy_hash` equals `capsule.policy_hash`
  - `statement.trace_root` equals `capsule.row_index_ref.commitment`
  - anchors match capsule anchor (if present)

Then it calls `zk_verify_geom(...., statement_hash=bytes.fromhex(statement_hash))`.

So the verifier is forcing the **same statement hash** to be used in the STARK verifier.

## (8) Optional Nova consistency check

If `nova_stats` exists, it checks:

- `nova_stats["nova_state"] == capsule["trace_commitment"]`  
Mismatch ⇒ E055\_NOVA\_STATE\_MISMATCH

## (9) Row index commitment verification

If row archive info exists:

- loads `chunk_roots` from inline/json/bin
- reconstructs k-ary Merkle root
- checks equals `row_index_ref.commitment`

So even before any DA sampling, it checks the row index root is consistent with the chunk roots list.

## (10) Authorship + ACL

- If ACL exists, authorship becomes required.
- It verifies recoverable secp256k1 signature recovers claimed pubkey.
- Then checks signer pubkey is authorized for the `policy_id` (status active).

## (11) DA audit (probabilistic sampling)

If `da_policy` is present:

- deterministically samples indices using a seed derived from:
  - `capsule_hash`
  - `anchor_ref`
  - `policy_id, policy_version`
- fetches sampled chunks from provider
- recomputes each sampled chunk root and verifies its Merkle inclusion proof up to the row root

This is why you call it “probabilistic\_da\_sampling.”

---

## What I think — where this stands overall

Based purely on what's in the zip:

**This is strong as a `capsule contract`**

You've got a coherent object model:

- **TraceSpec** = “what does this trace mean?”
- **Statement** = “what claim are we proving, bound to policy/root/inputs/anchors?”
- **Capsule hash** = “this entire bundle is integrity-protected”
- **Proof payload hashes** = “artifacts are content-addressed too”
- **Policy registry inclusion** = “governance pinning is possible”
- **Authorship/ACL** = “who’s allowed to produce capsules under this policy”
- **DA sampling** = “retrievability checks are enforceable”

That’s a *real* systems design. It’s not vibes.

## But it’s still a demo-shaped system

Because:

- Only one AIR (`GEOM_AIR_V1`) is concretely wired.
- The “adoption surface” for *other traces* is implied, not implemented in this snapshot.
- DA is implemented as **sampling + local archive provider**, not a full general DA layer story.

## My honest placement

- **“Proof object / verification contract” maturity:** high (you’re doing the right bindings and checks)
- **“General-purpose platform / others can adopt without you” maturity:** medium-low (Trace Adapter API is the missing headliner)
- **“Market wedge clarity” maturity:** medium (you’ve built the primitive; wedge depends on packaging and target)

If you want the cleanest one-liner for what this is **today**:

**“A policy-bound proof capsule format + verifier that prevents proof reuse under different semantics and supports deterministic DA sampling audits.”**

That’s already a *category*.

---

## The 3 biggest “what’s missing next” signals from the zip itself

These are not generic advice — they fall straight out of the code:

1. **Trace Adapter API (first-class)**

Right now, the pipeline hardcodes how to:

- define schema
- define TraceSpec
- define public inputs
- define Statement
- invoke prover/verifier

To make it adoptable, you need a standard interface like:

- `TraceAdapter.build_trace_spec()`
- `TraceAdapter.extract_public_inputs()`
- `TraceAdapter.statement_from_commitment(...)`
- `TraceAdapter.program_id / air_id / vm_id`

## 2. Policy semantics need teeth

You already hash the policy file and optionally pin it to a registry — great.

But “policy” will only matter if you define the **minimum required checks** that a verifier must enforce (beyond inclusion).

Right now policy is mostly “bound,” not “interpreted,” except via ACL and DA parameters.

## 3. DA story: name it precisely

Your implementation is: deterministic sampling of chunks from a known archive + Merkle proof checking.

That’s valuable, but it’s not “all DA problems solved.” It’s **audited retrievability** under a sampling policy.

That’s fine — it’s actually a wedge if you position it right.

---

If you want, I can do the next step in the same “look at the zip and be concrete” mode:

- I’ll write you the **exact wedge options that match what this system already enforces**, not speculative ones.
- And I’ll map each wedge to the *fields already in the capsule* (policy\_id, anchors, da\_policy, ACL), so the pitch sounds inevitable.

Alright. Here are wedge options that are **compatible with what your zip already enforces today** (TraceSpec/Statement binding, policy hash + optional registry, authorship/ACL, deterministic DA sampling, artifact hashing), plus exactly how you’d pitch + implement each **using the fields you already have**.

I’m going to be a little ruthless about “what you can honestly claim *now*” vs what would be overclaiming.

---

## Wedge 1 — Verifiable backtests / “proof-carrying alpha reports”

**Who buys first:** trading teams, quant pods, prop-ish shops, crypto funds, even internal risk/compliance.

**Pain:** “Your backtest is a PDF + trust me.” No provenance, no reproducibility, no audit trail, easy to cherry-pick.

**What you deliver:** a capsule that proves:

- the *exact* backtest code/program variant ran (TraceSpec ties to schema/program id)
- on the *exact* dataset snapshot (anchor)
- under an agreed policy (policy hash + optional registry pinning)
- and the raw trace/dataset chunks are retrievable with a quantifiable sampling guarantee (da\_policy)

**How your current capsule fields map**

- **anchors**: use `dataset_hash`, `data_vendor_attestation`, `universe_hash`, `time_range`, or even “exchange snapshot hash”
- **policy\_hash**: encodes “no lookahead, slippage model, fees model, train/test split rules, max leverage, etc.”
- **policy\_id / policy\_version**: “strategy\_eval\_policy/v3”
- **public\_inputs**: final metrics (CAGR, Sharpe, max drawdown, turnover) + key intermediate checkpoints
- **authorship + ACL**: only approved “research runners” can produce capsules that count
- **da\_policy.verification\_level="probabilistic\_da\_sampling"**: “audited retrievability” of the raw trace or row archive

**Why this is a good wedge for your current system**

You don’t need “L1 rollup DA.” You need “audited reproducibility.” Your DA sampling is *already enough* to say: “if you challenge me, I can produce the underlying rows/chunks, and you can verify inclusion to the committed root.”

**Skeptic pushback you’ll get (and how to answer)**

- “Sampling isn’t full availability.”  Correct — don’t oversell. Call it **retrievability audit** with configurable ( $\delta$ ,  $\epsilon$ ) policy.
- “What if you cherry-pick a dataset?” That’s exactly why **anchors** exists. Make dataset anchoring non-optional in the policy.

---

## Wedge 2 — zkML evaluation receipts (model claims that can't be faked)

**Who buys first:** ML infra teams, model marketplaces, “model provider” startups, AI safety/compliance groups.

**Pain:** benchmark claims are easy to lie about; provenance is weak; dataset leakage is rampant.

**What you deliver:** “This model achieved X on dataset Y under eval policy Z, and here's a cryptographic artifact that binds it.”

### Field mapping

- `anchors`: `dataset_hash`, `benchmark_suite_version`, `promptset_hash`, `model_weights_hash`
- `policy_hash`: evaluation policy (no internet, deterministic seeds, allowed preprocessing, max context length, etc.)
- `public_inputs`: metrics + confusion matrices / aggregate stats
- `TraceSpecV1`: schema reference for trace rows (e.g., per-example logits hashes / loss summaries)
- `DA sampling`: lets an auditor spot-check example-level trace chunks without you dumping everything publicly

### Why it fits

Your design is basically a “receipt object” with stronger policy semantics than typical zk receipts.

### Hard truth / constraint

This wedge is *only credible* if you can make the trace adapter story clean for “any evaluation pipeline.” Right now zip is geom-shaped. But this wedge doesn’t require new cryptography — it requires better adapter UX.

---

## Wedge 3 — “Policy-bound compute attestations” for regulated workflows

**Who buys first:** fintech, enterprise, gov-ish orgs, internal audit teams.

**Pain:** proving you followed process controls (“we ran the approved pipeline on approved data”) is expensive and mostly paperwork.

**What you deliver:** a capsule that *is* the audit artifact: hash-bound policy + authorized signer + anchored inputs + retrievability checks.

### Field mapping

- `policy_registry_root + inclusion_proof`: this becomes huge here — it's "the policy was approved by governance."
- `authorship/ACL`: policy-bound signing authority ("only these keys can attest")
- `anchors`: change-ticket id, build artifact hash, config hash, dataset snapshot hash
- `statement_hash` in transcript: stops replaying an old proof under a new "approved" policy

### Why it's real

Unlike "build a rollup," this wedge sells into people who already pay for audit/compliance, and you give them cryptographic receipts instead of screenshots.

### Skeptic pushback

- "We can do this with Sigstore / attestations." Yes — so the differentiation must be: **you're attesting to the computation's correctness, not just that someone signed a build.** Your statement/proof binding is the value.
- 

## Wedge 4 — “Capsules as verifiable research artifacts” for multi-party collaboration

**Who buys first:** any team collaborating across trust boundaries (joint research, open competitions, bounty-based R&D, distributed teams).

**Pain:** people don't trust each other's experiment claims; reproducibility is painful.

**What you deliver:** a capsule as the unit of collaboration:

- one object to share
- deterministic verification
- stable error codes
- can be pinned to a registry or signed

### Field mapping

- `artifact_manifest.json`: content-addressed index of "what files does this capsule refer to"
- `anchors`: commit hash, repo tag, docker image hash
- `policy_hash`: "how experiments must be run"

- `authorship`: identities

#### **Why it's a good “fast demo wedge”**

You can ship this without integrating into chains or DA networks. It's a software distribution wedge.

---

## **Wedge 5 — “Bridge / oracle-grade provenance layer” (careful: only if you narrow scope)**

**Who buys first:** oracle teams, bridge teams, on-chain protocols that need trustworthy off-chain computation summaries.

**Pain:** on-chain contracts depend on off-chain processes (indexing, pricing, risk) with weak attestations.

#### **What you can honestly do now**

You can produce a capsule proving “off-chain computation X produced value Y under policy Z and anchored to block B.”

#### **Field mapping**

- `anchors`: L1 block hash / finalized checkpoint
- `public_inputs`: the oracle values being posted on-chain
- `policy_hash`: exactly what computation is allowed
- `authorship/ACL`: only authorized operators can produce capsules accepted by the contract

#### **What you should NOT overclaim yet**

This does not automatically become a decentralized oracle network. It becomes **auditable operator attestation** with cryptographic correctness and retrievability audit.

---

## **Wedge 6 — “Proof-carrying benchmarks” for ZK / systems performance**

**Who buys first:** ZK infra companies, provers, hardware teams, researchers.

**Pain:** performance claims (“we proved X in Y ms”) are unverifiable and cherry-picked.

## What you deliver

A capsule whose `public_inputs` include benchmark results, and whose trace proves the benchmark workload ran as specified.

## Field mapping

- `policy_hash`: benchmark rules
- `TraceSpec`: schema of performance trace
- `anchors`: hardware spec hash / environment hash

This is a surprisingly good wedge if you want credibility fast in ZK circles — but it's not usually a big business.

---

# Which wedge is “best” given what you have right now?

If we’re being brutally practical:

## Best immediate wedge (fastest to product):

### Wedge 1 (Verifiable backtests) + Wedge 3 (Policy-bound attestations)

Because:

- your policy registry + ACL + statement binding is *already* tailor-made for “compliance-like” claims
- “DA sampling” is plenty for **audited retrievability** (spot checks)
- you don’t need to solve global DA or build a chain

## Most “platform-y” wedge (bigger upside, slower):

### Wedge 2 (zkML eval receipts)

Huge market narrative, but requires the Trace Adapter API to be clean so others can plug in pipelines.

---

# How to pitch the wedge in one sentence (per option)

- **Backtests:** “A cryptographic backtest receipt that binds dataset snapshot + evaluation policy + metrics, and supports spot-checkable retrievability of the underlying run.”
  - **zkML eval:** “Verifiable benchmark receipts: model claims you can audit, tied to an exact dataset + eval policy.”
  - **Compliance:** “A policy-bound computation attestation that’s governance-pinnable and signer-restricted.”
  - **Oracle/bridge:** “Correctness proofs for off-chain compute outputs, anchored to chain state and bound to operator authorization.”
- 

## Two concrete GTM experiments you can run *without adding new crypto*

### Experiment A — “Challengeable backtest capsule”

1. Pick 1 strategy + 1 dataset snapshot.
2. Write a policy file that forbids lookahead/leakage and requires dataset anchoring.
3. Emit capsule with:
  - `anchors: dataset_hash + time range + vendor`
  - `public_inputs: key metrics`
  - `da_policy` configured to a strong-sounding but honest epsilon/delta
4. Invite someone to challenge: “Ask for chunk i, we’ll produce it, you verify inclusion.”

That demo sells your system better than any theory.

### Experiment B — “Governed policy registry”

1. Create a tiny policy registry root (Merkle over allowed policies).
2. Ship capsules that include an inclusion proof to that root.
3. Now you can say: “Only governance-approved evaluation policies are valid.”

That’s the difference between “cool proof” and “enterprise/protocol-ready artifact.”

---

# One last “don’t fool yourselves” note

Your system’s sharp edge is **semantic binding + governance + auditability**.

If you try to wedge as “data availability network” or “rollup competitor” too early, people will compare you to Celestia/EigenDA/etc. and you’ll lose on category mismatch. But if you wedge as **policy-bound proof capsules for off-chain computation**, your current design looks *purpose-built*.

If you tell me which world you want to win first—**finance backtests, zkML evals, or compliance attestations**—I’ll write:

- the exact capsule fields you should require in policy,
- the demo storyline,
- and the 5-slide wedge pitch (no fluff, just inevitability).

Wedge 6 is actually *sneaky good* — not as a “huge revenue day 1” product, but as a **credibility + adoption Trojan horse** that perfectly matches what your capsule system already does well.

## The honest premise

Right now, zk performance “benchmarks” are mostly:

- blog posts by teams with incentives (RISC Zero vs SP1, etc.) ([RISC Zero](#))
- some community attempts at standardization (ZK Bench / zk-Harness / zkvm-perf) ([ZK Bench](#))
- occasional broader meta-reports (e.g., Fenbushi’s 2025 benchmarking overview) ([Fenbushi](#))

The missing thing is: **a result format that’s cryptographically bound to “what ran,” under “what rules,” on “what inputs,” with “what artifacts,” and optionally “who is authorized to claim it.”**

That is *literally* what your capsule schema is.

---

## What wedge 6 becomes in your world

**“Proof-carrying benchmarks”** = a benchmark result that is only accepted if it comes with a capsule that verifies.

Not “trust me bro, we ran it.”

More like: “Here’s the exact workload definition + policy + inputs + proof + artifact hashes + audit hooks.”

This matches the benchmarking standards vibe the ZKProof community has talked about for years: benchmark frameworks and recommended practices, plus qualitative properties/assumptions that need to be declared. ([ZKProof](#))

---

## Why wedge 6 fits your zip today

Your current capsule/verifier machinery already supports the essential ingredients:

1. **Workload identity (TraceSpecV1)**
    - `trace_format_id, record_schema_ref, encoding_id, field_modulus_id`
    - This is your “benchmark definition fingerprint.”
  2. **Claim binding (StatementV1 in transcript)**
    - binds trace root + public outputs + anchors + policy hash
    - stops re-labeling the same proof as a different benchmark.
  3. **Rules of the benchmark (policy hash + optional registry inclusion)**
    - you can publish “Benchmark Policy vN” and pin it via a registry root.
    - now people can’t silently change methodology.
  4. **Artifact integrity**
    - proof payload hashes let you content-address the proof files.
  5. **DA sampling**
    - for benchmarks, this is *perfect* as “auditability”: you can spot-check logs/chunks without demanding everyone publish massive traces.
  6. **Authorship + ACL**
    - lets you do “certified runs”: only keys in the policy ACL can publish “official” benchmark capsules.
- 

## The key limitation (and how to position it honestly)

You **cannot** cryptographically prove wall-clock time in a trustless way from pure software proofs alone.

So don’t sell “we prove runtime.” Sell:

### What you can prove

- the computation/workload was executed correctly
- under declared constraints (policy)
- with declared inputs/versions (anchors)
- and the artifact set is integrity-protected

- plus you can include measurements as *signed telemetry* and make them auditable/reproducible.

## How the ecosystem already handles this

Even existing benchmark frameworks emphasize how many “factors affect performance” and the difficulty of direct comparisons. ([ZK Bench](#))

And tools like zkvm-perf are basically trying to standardize the process around running benchmarks. ([GitHub](#))

Your wedge is: **standardize the process + make it cryptographically enforceable.**

---

## The clean “product shape” for wedge 6

If you do wedge 6, the product is basically:

### 1) A benchmark suite spec

- “workloads” = named TraceSpecs (and/or “program IDs”)
- “policies” = methodology: compiler flags, CPU pinning, dataset sizes, warmups, GPU model constraints, etc.
- “required anchors” = git commit hash, docker image digest, prover version, hardware manifest

### 2) A runner that emits capsules

- wraps existing tools (zkvm-perf, zk-bench style harnesses) but outputs **your capsule** at the end ([ZK Bench](#))

### 3) A verifier + leaderboard that only accepts verified capsules

- no capsule, no leaderboard entry
- “certified track” requires ACL signing

### 4) Optional: “certified lab” program

- keys represent labs/hardware providers
  - ACL restricts who can publish certified results
- 

## How to pick workloads so this wedge lands

You want workloads that match how teams already publish comparisons:

- “real-world workload” style comparisons (SP1 and RISC Zero both benchmark specific workloads and blocks) ([Succinct](#))
- “microbenchmarks” (hashes, FFT/MSM-ish components, recursion steps, etc.)
- “proof system stack” benchmarks (zk-Harness explicitly aims to cover the stack) ([GitHub](#))

The Fenbushi benchmarking article suggests the market is actively paying attention to cross-zkVM performance comparisons right now. ([Fenbushi](#))

So you’d be riding an existing wave, not inventing demand.

---

## Where wedge 6 stands commercially

This is the real trade:

### What wedge 6 is *great* for

- instant credibility in the ZK ecosystem
- forcing a standard “capsule interface” into people’s workflows
- becoming the neutral layer everyone references

### What wedge 6 is *not* great for (alone)

- big revenue early, unless you attach it to:
  - certification
  - enterprise procurement (help teams choose provers/hardware)
  - “benchmarking as a service” for teams shipping prover networks / rollups

---

## If you want the tightest possible pitch for wedge 6

“We built a cryptographically verifiable benchmark artifact format. Benchmarks become receipts: workload definition + policy + input anchoring + proof + artifact hashes + optional certified signer. No more unverifiable performance claims.”

...and the subtext is: *once people adopt capsules for benchmarks, adopting capsules for everything else becomes natural.*

---

# Concrete next step (no new cryptography required)

If you say “yes wedge 6,” the fastest win is:

1. Ship **capsule-bench** CLI:
  - runs a benchmark workload (even if initially only your GEOM demo + 1–2 microbenchmarks)
  - emits capsule + artifacts
  - verifies locally
2. Ship a tiny website that:
  - uploads capsule
  - verifies capsule
  - posts result to a leaderboard
3. Add a **policy registry root** and a **certified track**:
  - community track: anyone can post
  - certified track: must be signed by keys in ACL

That’s enough to be “a thing” people can point at.

---

If you want, tell me which direction you want wedge 6 to be:

1. **ZKVM leaderboard (neutral ecosystem credibility)**
2. **Enterprise “certified benchmark reports” (procurement/compliance angle)**
3. **Prover-network sales enablement (help teams market performance with receipts)**

...and I’ll write the exact capsule field requirements + the policy schema + the demo storyline you should run with.

Yeah — **1 + 3 is the right combo because they reinforce each other.**

- **(1) Neutral leaderboard** gives you legitimacy + attention + a standard everyone references.
- **(3) Sales enablement for prover networks / zkVM teams** turns that legitimacy into distribution + money (teams *want* to market performance claims, but the space is noisy and trust is low).

And the best part: this is already where the ecosystem is moving. There are multiple active benchmark efforts (tooling + repos + reports), but very few “**cryptographically-bound benchmark receipts**” as the unit of truth:

- Tooling exists (e.g., **zkvm-perf**) ([GitHub](#))
- Other benchmark harnesses exist (**zkvm-bench**, **zk-Harness**) ([GitHub](#))

- People are explicitly trying to do standardized comparisons (Fenbushi's 2025 report; Ethereum Foundation zkEVM benchmarking post) ([Fenbushi](#))

## The wedge statement

**"We make zkVM benchmarks receipt-based."**

No receipt (capsule) → it doesn't count. And "receipt" doesn't just mean logs; it means **(TraceSpec + Statement + policy hash + artifact hashes + optional signer/ACL + audit hooks)**.

That's exactly what your capsule design already supports.

---

## Product shape for (1) Neutral leaderboard

### What you ship

A public leaderboard that only accepts verified capsules.

### Tracks

1. **Community track**: anyone can submit a capsule, must verify.
2. **Certified track**: capsule must include `authorship` and pass `ACL` checks for a given `policy_id` (your verifier already supports this pattern).

### Why this works

Most benchmarking repos can *run tests*; your differentiator is a **verification gate** and a portable artifact format.

### Bench workloads

You don't need 50 benchmarks. You need **5–8** that are:

- widely understandable,
- hard to game,
- representative (CPU-heavy, GPU-heavy, recursion-ish, memory-ish).

You can even *wrap* existing benchmark suites rather than inventing everything:

- zkvm-perf exists specifically as a benchmarking tool for zkVMs ([GitHub](#))
- zkvm-bench provides an "apples-to-apples" comparison setup across multiple zkVMs ([GitHub](#))

Your move is: **output a capsule at the end.**

---

## Product shape for (3) Sales enablement (the money printer)

Here's what zkVM/prover-network teams want *badly*:

- "Here are our numbers."
- "They're reproducible."
- "They're fair."
- "Here's the raw methodology."
- "Here's a badge / proof it's legit."

This is already playing out in the wild: teams publish benchmark claims (SP1 posts benchmark articles; RISC Zero posts "zkVM 1.0" performance claims) ([Succinct](#)) — but readers always ask "under what settings, what hardware, what version, what shortcuts?"

### What you sell

Verified Benchmark Packs:

- a capsule bundle + signed certification + a shareable "Benchmark Receipt" page
- optional "Run in our lab" service (you control hardware + policy registry)
- "Competitive comparison report" (capsules for multiple systems under identical policy)

Fenbushi's 2025 benchmarking report is basically evidence that *market attention exists* for standardized zkVM comparisons ([Fenbushi](#)). You're productizing that into receipts.

---

## How your *existing* capsule fields become the standard (no handwaving)

To make this wedge crisp, define these as **required** in the benchmark policy (so verifiers enforce them):

### Required anchors (anti-gaming)

- `git_commit` (benchmark runner + workload)
- `docker_image_digest` (or build hash)
- `hardware_manifest_hash` (CPU/GPU model, RAM, OS/kernel, driver versions)
- `workload_input_hash` (fixed input set)

- optionally `cloud_instance_type` (when running in cloud)

## Required public inputs (what shows up on leaderboard)

- proof size
- prover time as *declared measurement* (not cryptographically proven, but policy-governed)
- peak memory
- cycles/steps (workload size parameter)
- verification time

## Required policy fields

- warmup rules
- repetitions + aggregation rule (median/mean)
- allowed flags (or explicitly disallow compiler accel/custom params unless stated)
- whether GPU is allowed, which GPU class
- whether recursion is included

Your verifier already binds `policy_hash` into the statement hash, so the proof can't be "relabelled" under a different policy later.

---

## The “credibility lock”: Policy registry + Certified Keys

This is how you prevent the leaderboard from becoming a circus:

- Publish a **policy registry root** (Merkle root of approved policies).
- Submissions must include inclusion proofs for “Certified” track.
- Certified runs require signature + ACL authorization per policy.

This is exactly the kind of governance/provenance story serious benchmarking needs, and it maps to what your verifier already supports.

---

## Practical rollout plan (fast)

### v0 (2 weeks energy)

- 3 workloads
- community track only

- upload capsule → verify → show results

## v1 (adds the wedge “teeth”)

- policy registry + certified track
- fixed runner repo + docker image
- “compare two zkVMs under identical policy” template

## v2 (turn into distribution + revenue)

- “Verified Benchmark Pack” for zkVM teams (they embed your badge)
  - optional hosted runs (your hardware)
  - enterprise/protocol-facing reports (“pick a prover network”)
- 

## The one pitfall to avoid

Don’t claim you “prove wall-clock time.”

What you *can* claim is stronger and cleaner:

**“We make benchmarks reproducible and policy-bound, with cryptographic integrity over the artifacts and a verifier-gated submission format.”**

That’s a very defensible line, and it’s aligned with the direction of current benchmarking efforts (standardized harnesses + cross-system comparisons) ([GitHub](#)).

---

If you say “go,” I’ll write you:

1. a **Benchmark Policy v1 JSON schema** (the exact fields + rules),
2. the **capsule submission contract** (what a run must include),
3. the **leaderboard categories** (so it doesn’t devolve into apples vs oranges).

Yeah — **1 + 3 is the cleanest “inevitable” wedge**, and you already have the right primitives in your capsule system to make it real. Let me lay out the *actual* product thesis in a way that’s hard to argue with.

## The core insight

Benchmarking in ZK is messy for two reasons:

1. **Comparability is structurally hard** (too many knobs; comparisons can unintentionally favor one system). zk-Harness literally starts from this premise: lots of zkSNARK libraries exist, but “various factors make it difficult to compare different proof systems without favoring some approaches over others.” ([GitHub](#))
2. The ecosystem *increasingly needs benchmarking as infrastructure*, not as blog posts. The Ethereum Foundation’s zkEVM team is explicitly benchmarking zkVMs to evaluate **real-time proving** and worst-case performance, because failing RTP threatens liveness/finality (their framing, not yours). ([zkEVM](#))

So the wedge isn’t “we benchmark zkVMs.” Lots of repos do. (zkvm-perf, brevis’ zkvm-bench, a16z’s repo, etc.) ([GitHub](#))

The wedge is:

## Benchmarks become “receipts.”

A benchmark result is only accepted if it comes with a **verifiable capsule** that binds:

- *what ran* (TraceSpec / program identity)
- *under what rules* (policy hash + optional policy registry inclusion)
- *on what inputs* (anchors)
- *with what artifacts* (proof + payload hashes)
- *who is allowed to claim “certified”* (authorship + ACL)
- *with audit hooks* (your deterministic sampling retrievability check)

This turns benchmarking from “content” into **a protocol**.

---

## Why 1 + 3 works as a pair

### (1) Neutral leaderboard = legitimacy + standard adoption

If you become the place where results are posted **only if they verify**, you set the format.

### (3) Sales enablement = distribution + money

Teams *want* to publish performance claims, and they also *want* a credible way to say “fair apples-to-apples.” Brevis’ repo explicitly emphasizes same hardware, same Rust program, and no compiler acceleration/custom params to keep things fair. ([GitHub](#))

You can sell them **“Verified Benchmark Packs”** that:

- produce capsules under a strict policy
- show up on the neutral leaderboard
- come with a shareable “Receipt Page” + badge

The leaderboard makes the receipts matter; the teams fund the receipts and push them everywhere.

---

## The key: be honest about what's proven vs measured

Your capsule can't "prove wall-clock time" in a trustless way by itself. That's fine.

**What you *can* claim (and it's strong):**

- Correctness of the benchmark workload execution/proof (your ZK verifier)
- Integrity of declared measurements (policy-bound + signed + anchored)
- Reproducibility constraints (docker/git/hardware anchors)
- Auditability (spot-check retrievability via your sampling mechanism)

This is aligned with EF's viewpoint that benchmarking is about defining constraints/fixtures and iterating measurements in a loop—not one-off conclusions. ([zkEVM](#))

---

## What you should actually build (minimum viable but "real")

### 1) A "Benchmark Policy v1" that verifiers enforce

Your *policy file* becomes the constitution. It should require:

**Required anchors (anti-gaming)**

- `benchmark_suite_id + suite_version`
- `workload_id + workload_input_hash`
- `git_commit` of the runner/workload
- `docker_image_digest` (or build hash)
- `hardware_manifest_hash` (CPU/GPU model + RAM + OS/kernel + driver versions)
- optional: `cloud_instance_type` (Brevis uses a specific instance type in their apples-to-apples setup; policies can standardize that.) ([GitHub](#))

**Required measurement procedure**

- warmups
- repetitions
- aggregation rule (median, trimmed mean)

- fixed flags allowed / disallowed (e.g., “no compiler acceleration” track vs “optimized” track)

### Required public outputs

- prove time, peak RSS, proof size, verify time, and the workload size parameter(s)

You already have the mechanism to bind the policy hash into the *statement hash in the transcript* (so results can’t be relabeled under a different methodology afterward).

## 2) A runner CLI that emits capsules (wrapping existing ecosystems)

You don’t need to fight existing tools. You *wrap them* and emit your receipt at the end.

- `zkvm-perf` already positions itself as a benchmarking tool for zkVM implementations with automated and manual modes. ([GitHub](#))
- brevis’ `zkvm-bench` gives an explicit “apples-to-apples” template (same hardware, same Rust program, no accel). ([GitHub](#))
- a16z’s repo is another community focal point for installing/running multiple zkVMs side-by-side. ([GitHub](#))

Your move: **standardize output into capsules** and make the verifier the gate.

## 3) The leaderboard (two tracks at first)

### Community track

- Must verify capsule + policy hash + anchors.

### Certified track

- Requires `authorship` signature and ACL allowlist for that `policy_id`.
- Optionally requires policy registry inclusion proof (so “Certified Policy v1” is pinned to a registry root).

That one design choice prevents the leaderboard from becoming a marketing spam wall.

---

## The leaderboard categories that avoid apples-vs-oranges

This is where most benchmark efforts fail. You need strict, obvious buckets:

1. **No-Accel / Baseline**

- matches the “same program, no compiler accel/custom params” ethos that brevis calls out. ([GitHub](#))
  - 2. **Optimized (declared knobs allowed)**
    - allowed acceleration, but all knobs must be declared in anchors + policy.
  - 3. **Recursion / aggregation track**
    - optional; only comparable if methodology is pinned.
  - 4. **Ethereum block proving track (longer-term)**
    - aligns with EF’s active benchmarking pipeline for zkVMs for Ethereum, including worst-case fixtures and RTP framing. ([zkEVM](#))
- 

## Why “now” is actually good timing

EF is actively building a benchmarking pipeline and explicitly cares about worst-case, reliability, and RTP constraints for mainnet readiness. ([zkEVM](#))

That means: the ecosystem is going to converge on benchmark standards anyway. The only question is whether the “unit of truth” is:

- spreadsheets + blog posts, or
- **verifiable receipts**

You have the ingredients to make it the second.

---

## If you want the sharpest one-sentence positioning

“We turn zkVM benchmarks into verifiable receipts—policy-bound, input-anchored, artifact-hashed, and optionally certified—so performance claims become reproducible objects, not marketing.”

If you want, next message I’ll draft:

- a concrete **Benchmark Policy v1 JSON** (fields + required anchors + tracks),
- a **capsule submission contract** (what your server checks before accepting),
- and the first **5 workloads** that are “high signal” and hard to game.

Yep. Here are **5 workloads** that (a) are *high-signal*, (b) map to real-world ZK/blockchain use, and (c) are hard to “game” without it being obvious. Then I’ll give you a **Benchmark Policy v1 JSON** + a **capsule submission contract** your server can enforce.

## The 5 workloads

## 1) Keccak-256 streaming throughput (fixed bytes, fixed schedule)

**Why it's high-signal:** Keccak is everywhere (Ethereum, Merkleization, transcripts). It stresses pure CPU + memory bandwidth in a very "honest" way.

### Workload definition

- Input sizes: 1 MiB, 16 MiB, 64 MiB
- Operation: `keccak256(buffer)` repeated  $R$  times with a chaining rule (next buffer = previous digest expanded deterministically) so you can't short-circuit
- Output (public): final digest + bytes processed

### Hard-to-game knob

- Inputs are *deterministically generated* from `suite_seed + size` (not user-provided)
- 

## 2) Merkle tree build + random openings (hash-heavy + random access)

**Why high-signal:** This matches your world (commitments, chunk roots). It also hits cache/memory patterns.

### Workload definition

- Leaves:  $N = 2^{20}$  leaves of 32 bytes (or  $2^{18}$  if too heavy initially)
- Hash: Keccak (or SHA-256 if you want non-EVM bias; you can have two variants)
- Task A: build root
- Task B: verify  $Q=1024$  inclusion proofs at random indices derived from seed
- Output (public): root + verification accumulator hash

### Hard-to-game knob

- Indices come from seed; proofs must be consistent with the built tree
- 

## 3) Signature verification batch (secp256k1 ECDSA verify)

**Why high-signal:** Big-int-ish arithmetic shows up; also extremely relevant to crypto systems.

### Workload definition

- Verify  $B=1024$  signatures over fixed messages

- Messages, pubkeys, signatures are deterministically derived from seed (or taken from a canonical corpus with a published hash)
- Output (public): count of valid signatures + final accumulator digest

#### Hard-to-game knob

- The dataset is pinned by `workload_input_hash` (anchor), so nobody swaps “easy” signatures
- 

### 4) Bytecode interpreter micro-EVM (control-flow + memory + hashing)

**Why high-signal:** This approximates “real computation” for zkVMs without requiring full Ethereum block proving.

#### Workload definition

- A tiny fixed instruction set: `ADD`, `MUL`, `XOR`, `SHL`, `SHR`, `MLOAD`, `MSTORE`, `JUMP`, `JUMPI`, `KECCAK_CHUNK`
- Execute  $S$  steps over a memory buffer with a deterministic program generated from seed
- Output (public): final state hash + step count

#### Hard-to-game knob

- Program is seeded and must match `program_hash` / `workload_input_hash`
- 

### 5) Memory bandwidth + branching stress (“STREAM + branchy loop”)

**Why high-signal:** It exposes VM overhead, memory model efficiency, and interpreter/JIT differences. Also reveals “cheating” quickly.

#### Workload definition

- STREAM triad style: `a[i] = b[i] + scalar*c[i]` over large arrays
- Plus a branchy pass: conditional swaps / thresholding using a seeded threshold vector
- Output (public): checksum hash of arrays + bytes touched

#### Hard-to-game knob

- Checksums are hashed, not just summed, to avoid algebraic shortcircuiting
-

# Benchmark Policy v1 JSON (concrete)

This is **the policy file** you hash and bind into the statement transcript. It's written to be *machine-checkable* by your verifier/server.

```
{
  "schema": "bef_benchmark_policy_v1",
  "policy_id": "bef_bench_policy_v1",
  "policy_version": "1.0.0",
  "suite": {
    "suite_id": "bef-zkvm-leaderboard",
    "suite_version": "2025-12",
    "suite_seed_hex": "9b3b7a2c1e4d...deadbeef",
    "workload_registry_hash": "sha256:REPLACE_WITH_HASH_OF_WORKLOADS_JSON"
  },
  "tracks": [
    {
      "track_id": "baseline_no_accel",
      "display_name": "Baseline (no accel)",
      "rules": {
        "forbid_custom_precompiles": true,
        "forbid_handwritten_asm": true,
        "forbid_jit": false,
        "forbid_gpu": true,
        "allowed_optimizations": ["-O2"],
        "forbidden_optimizations": ["profile-guided", "link-time-opt"],
        "require_deterministic_build": true
      }
    },
    {
      "track_id": "optimized_declared",
      "display_name": "Optimized (declared)",
      "rules": {
        "forbid_custom_precompiles": false,
        "forbid_handwritten_asm": false,
        "forbid_jit": false,
        "forbid_gpu": false,
        "require_full_knob_disclosure": true,
        "require_deterministic_build": true
      }
    }
  ],
}
```

```
"required_anchors": {
    "git_commit": true,
    "runner_repo": true,
    "docker_image_digest": true,
    "hardware_manifest_hash": true,
    "workload_input_hash": true,
    "toolchain_hash": true,
    "os_fingerprint_hash": true
},
"required_public_outputs": {
    "prover_time_ms": true,
    "peak_rss_mb": true,
    "proof_size_bytes": true,
    "verify_time_ms": true,
    "workload_size_params": true,
    "result_digest_hex": true
},
"measurement_procedure": {
    "warmup_runs": 3,
    "measured_runs": 10,
    "aggregation": "median",
    "clock_source": "monotonic",
    "pin_cpu_affinity": true,
    "disable_turbo_if_possible": true,
    "notes": "If turbo cannot be disabled, report it in hardware manifest."
},
"workloads": [
    {
        "workload_id": "keccak_stream_v1",
        "params": { "sizes_bytes": [1048576, 16777216, 67108864], "rounds": 4 }
    },
    {
        "workload_id": "merkle_build_open_v1",
        "params": { "leaves_pow2": 20, "openings": 1024, "hash": "keccak256" }
    },
    {
        "workload_id": "ecdsa_verify_batch_v1",
        "params": { "curve": "secp256k1", "batch": 1024 }
    }
]
```

```

    "workload_id": "micro_evm_interp_v1",
    "params": { "steps": 2000000, "mem_bytes": 16777216 }
},
{
    "workload_id": "mem_stream_branch_v1",
    "params": { "n_elems": 33554432, "passes": 3 }
}
],
"da_policy": {
    "verification_level": "probabilistic_da_sampling",
    "delta": 0.1,
    "epsilon": 1e-6,
    "min_k_samples": 64
},
"certification": {
    "require_signature_for_certified_track": true,
    "acl_schema": "bef_acl_v1",
    "policy_registry": {
        "enabled": true,
        "registry_root_hash_type": "sha256",
        "require_inclusion_proof_for_certified_track": true
    }
}
}

```

## Companion anchor formats (what the capsule must include)

Your capsule should include these keys under `anchors` (or a dedicated `benchmark` block—either is fine as long as the statement binds it):

- `git_commit`: commit of runner/workload repo
- `runner_repo`: canonical repo identifier
- `docker_image_digest`: immutable image digest
- `hardware_manifest_hash`: sha256 of a JSON hardware manifest
- `toolchain_hash`: sha256 of toolchain manifest (rustc/cargo/clang versions, etc.)
- `os_fingerprint_hash`: sha256 of OS/kernel/driver manifest
- `workload_input_hash`: sha256 of generated inputs (or canonical dataset tarball digest)
- `track_id`: one of the policy's tracks
- `workload_id`: one of the policy's workloads

---

# Capsule submission contract (server-side checks)

This is what your leaderboard server enforces **before** it accepts a run.

## A) Structural + integrity checks

1. **Parse + schema**
  - capsule must be a dict
  - `schema == "bef_capsule_v1"`
2. **Canonical hash**
  - recompute capsule hash (after removing `capsule_hash`)
  - must equal `capsule.capsule_hash`
3. **TraceSpec binding**
  - must contain `trace_spec + trace_spec_hash`
  - recompute and match
4. **Statement binding**
  - must contain `statement + statement_hash`
  - recompute and match
  - statement must bind:
    - `trace_spec_hash`
    - `policy_hash`
    - `trace_root` (row commitment root)
    - anchors (or benchmark block)
    - public outputs
5. **Proof artifact integrity**
  - if capsule includes `sha256_payload_hash` for proof formats, recompute and match
  - proof must verify under the verifier using the provided statement hash

## B) Policy + track enforcement

6. **Policy hash match**
  - server computes `sha256(policy_file_bytes)` and must equal capsule `policy_hash`
7. **Certified track gating**
  - if submission is for `certified`:

- require policy registry inclusion proof against a server-configured `policy_registry_root`
  - require `authorship` signature
- 8. **ACL enforcement (certified)**
  - require ACL file (or server-known ACL registry)
  - signer pubkey must be authorized for this `policy_id`
- 9. **Track rule checks**
  - enforce track constraints:
    - baseline track forbids GPU + custom precompiles + etc.
    - optimized track requires knob disclosure

These are enforced by requiring anchors/manifests that declare them + rejecting missing fields.

## C) Anchor requirements (anti-gaming / reproducibility)

- 10. **Required anchors present**
  - must contain all keys listed in `required_anchors`
  - `workload_id` and `track_id` must be in policy
- 11. **Hardware manifest hash matches**
  - user uploads `hardware_manifest.json`
  - server hashes it and matches `hardware_manifest_hash`
- 12. **Toolchain + OS fingerprint hashes match**
  - same idea: require uploaded manifests, hash them, match anchors
- 13. **Workload input hash matches**
  - either:
    - inputs are deterministically generated from policy seed → server recomputes expected hash, OR
    - inputs are uploaded as a blob/tarball → server hashes and matches `workload_input_hash`

## D) Output requirements (leaderboard fields)

- 14. **Required public outputs present**
  - must contain all metrics keys required by policy
  - check types/ranges (e.g., non-negative, reasonable bounds)
- 15. **Consistency checks**
  - `proof_size_bytes` must match actual proof file size
  - `verify_time_ms` must be from a declared verifier version (anchored)
  - `workload_size_params` must match policy workload params

## E) DA / retrievability audit (what you already support)

## 16. DA policy sanity

- must include `verification_level="probabilistic_da_sampling"`
- enforce `k_samples >= min_k_samples`

## 17. Server-performed sampling

- server derives audit seed deterministically (from capsule hash + anchors per your design)
  - server requests sampled chunks from the declared provider (or from uploaded row archive)
  - server verifies Merkle inclusion up to `trace_root`
  - failures reject the submission (or mark as “DA failed”)
- 

# Recommended “day 1” simplification (still honest)

To ship quickly without weakening integrity:

- Require submissions to upload a **bundle**:
  - capsule
  - proof artifact(s)
  - row archive (or chunk roots + sampled chunks)
  - hardware/toolchain/os manifests
- Server stores those, then runs verification + sampling.

You can later support remote DA providers.

---

If you want, next I can also generate:

- a `workloads.json` registry file (the canonical definition you hash into `workload_registry_hash`)
- example `hardware_manifest.json`, `toolchain_manifest.json`, and `os_fingerprint.json`
- a “badge rules” spec: when a run gets  Verified /  Certified /  DA-Audited

Yeah — **the pipeline can (and should) automatically gather + shape the whole submission UX**. In fact, *your current geom pipeline already does ~60–70% of it* (it computes hashes, writes the row archive, emits capsule + proof artifacts). What’s missing is basically: **(a) auto-manifests + (b) bundling + (c) a clean adapter interface so it’s not geom-only**.

Here's the clean breakdown: what the system can auto-do vs what an adopter has to do.

---

## Can the pipeline form the “UX shape” automatically?

### ✓ Yes, for everything that is “machine-observable”

Your runner can automatically:

- compute `policy_hash`
- compute `trace_spec_hash` and `statement_hash`
- compute `capsule_hash`
- compute payload hashes for proof files
- export `row_archive/` + chunk roots + metadata
- run local verification
- produce a single **bundle** (capsulepack) ready to upload

And it can automatically **collect**:

- hardware fingerprint (CPU/GPU/RAM)
- OS/kernel + driver versions
- toolchain versions (rustc/cargo, clang, etc.)
- proof sizes
- verification time (measured)
- prover time + peak RSS (measured)

None of that requires humans.

### ⚠ What it cannot “magically know”

It can't infer:

- what *your* computation semantics are (that's the Trace Adapter's job)
- what anchors are “the right ones” for your domain (policy decides this)
- whether you “used forbidden optimizations” unless you require disclosure and/or run in a controlled environment (docker + pinned runner)

So you don't prove “no JIT.” You enforce it by:

- (baseline track) running inside a standardized container/harness you provide, and
  - requiring manifests + disclosure + reproducible build.
-

# What does an adopter have to do?

This depends on what “adopter” means: **(A) someone submitting a benchmark run vs (B) a team integrating a new VM/computation into the capsule framework.**

## A) A *benchmark submitter* (most users)

If you ship `capsule-bench`, the submitter does basically:

1. **Install runner**

```
cargo install capsule-bench # or curl | sh
```

2. **Run**

```
capsule-bench run --workload merkle_build_open_v1 --track baseline_no_accel
```

3. **Submit**

```
capsule-bench submit --pack out/capsulepack.tgz
```

That's it. They never hand-write hashes, never craft manifests.

If they want **Certified track**, add:

```
capsule-bench run ... --sign --key ~/.bef/keys/secp256k1.pem
```

## B) A *system integrator* (zkVM/prover team adding support)

This is the real “adoption surface” question. For them, the work is:

### 1) Implement a Trace Adapter (the only real integration cost)

They provide a small interface like:

- `trace_format_id` (e.g. `SP1_VM_V1`, `RISC0_VM_V2`, `MYVM_AIR_V1`)
- `record_schema_ref` (hash of the row schema they output)
- `encode_trace(rows) -> stc_rows` (or whatever your STC ingestion format is)
- `extract_public_outputs(run) -> dict`
- `anchors_required(run) -> dict` (git commit, docker digest, workload input hash, etc.)

- optional `da_export(trace) -> row_archive`

Everything else (hashing, statement building, capsule emit, bundle) is your framework.

## 2) Plug into your runner wrapper

They essentially implement:

```
capsule-bench run --backend myvm
```

## 3) (Optional) Provide a verification artifact path

So your verifier can call:

- `zk_verify_myvm(proof, statement_hash, vk, ...)`

That's it. They don't have to understand your capsule internals beyond a schema.

---

# The “capsulepack” UX shape you should standardize

Make the output of **every run** identical:

```

capsulepack/
  capsule.json
  policy.json      (or policy_id + inclusion proof)
  proofs/
    geom.proof.bin
    geom.vk.bin     (if needed)
  row_archive/
    chunk_roots.bin
    chunk_roots_digest.txt
    ... optionally sampled chunks for DA
  manifests/
    hardware.json
    os.json
    toolchain.json
  measurements.json (prover_time_ms, peak_rss_mb, verify_time_ms, etc.)
  logs/
    run.log

```

A submitter just uploads **one tarball**. Your server does:

- unpack
  - verify capsule hash + statement/trace spec hash
  - verify proof
  - verify required anchors/manifests
  - do DA sampling audit (either from uploaded archive or remote provider)
- 

## So: what is the adopter *actually* responsible for?

### If they're just submitting

- run one command
- upload one file
- optionally sign with a key

### If they're integrating a new backend/computation

They are responsible for **exactly one thing**:

Provide a deterministic way to produce a trace + public outputs + anchors in a known schema (Trace Adapter), so your framework can bind and verify it.

Everything else should be framework-provided.

---

If you want, I can write the **exact TraceAdapter interface** as a minimal Python/TypeScript spec (like 15–25 lines of required methods + expected outputs), and the **capsulepack schema** + the server-side validation order — so “adopter effort” becomes a checklist you can hand to any zkVM team.